

Generative Adversarial Transformers

John Salvador

Sri Ram Pavan Kumar Guttikonda

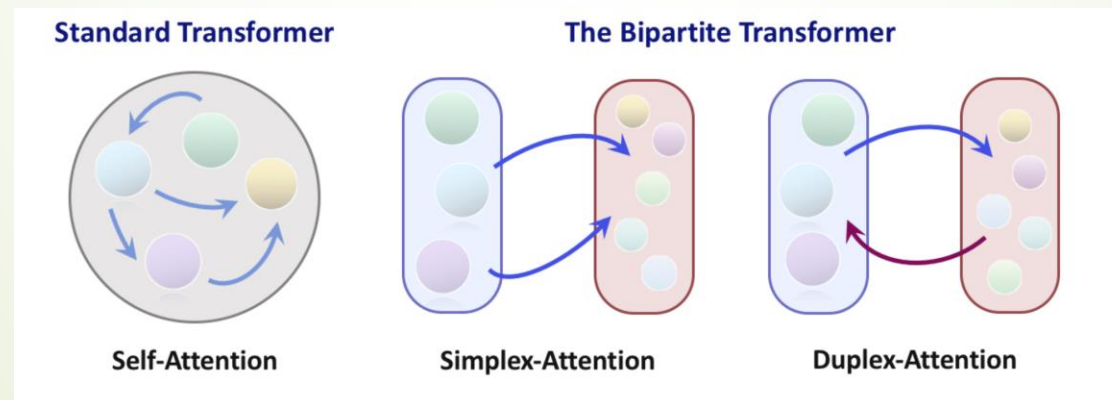


Contents

- **Contribution**
- Inspirations
- Attention
 - Vanilla Attention
 - Bipartite Attention
 - Simplex Attention
 - Duplex Attention
- Testing Results
 - FID Scores
 - Generated Images
- Code Overview
 - Environment Setup
 - Quick-start commands
 - Issues

Contributions

- Authored by Drew A. Hudson and C. Lawrence Zitnick
- Unifies GANs and Transformers into a single architecture
 - Allows different components to contribute to image generation
- Introduction of Bipartite attention for bilinear efficiency $O(mn)$
 - Better performance scaling for higher resolutions



Contents

- Contribution
- **Inspirations**
- Attention
 - Vanilla Attention
 - Bipartite Attention
 - Simplex Attention
 - Duplex Attention
- Testing Results
 - FID Scores
 - Generated Images
- Code Overview
 - Environment Setup
 - Quick-start commands
 - Issues

Inspiration From Human Perception

- the **bottom up** processing, proceeding from the retina up to the cortex, as local elements and salient stimuli hierarchically group together to form the whole.
- **Top-down** processing: Background knowledge is used to make inference





Image Source: Emily Roberts <https://www.verywellmind.com/bottom-up-processing-and-perception-4584296>
Ronald Sahyouni

Preceding Works

- The traditional convolution networks does not reflect this bidirectional nature that so characterizes the human visual system.
- Hard to produce diverse images in GAN
- The Transformer complexity for attention is $O(n^2)$.
- The Style GAN provides no means to control the style of a localized regions within the generated image.

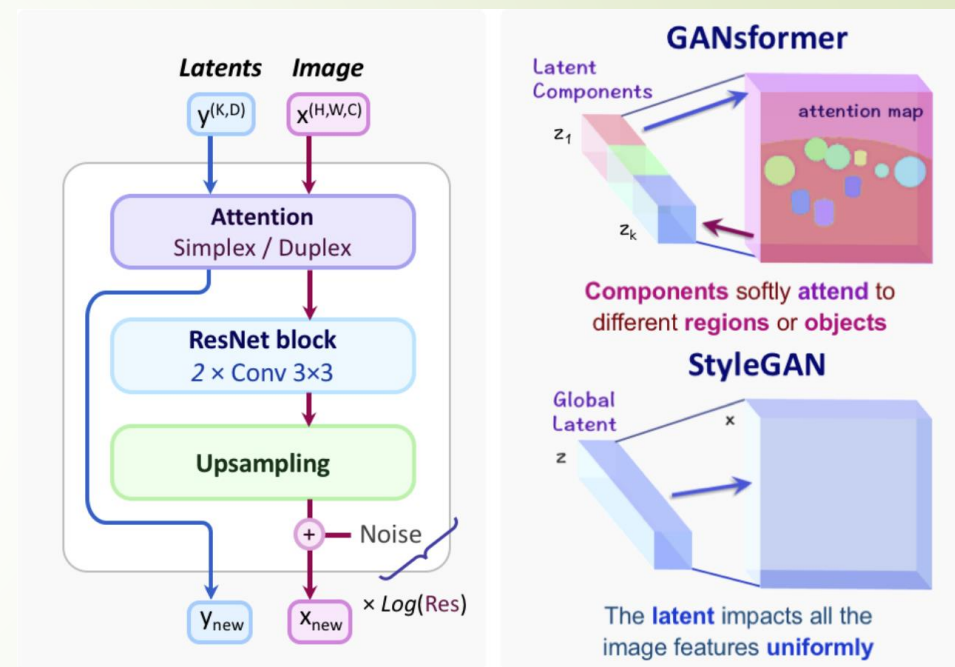


Contents

- Contribution
- Inspirations
- **Attention**
 - Vanilla Attention
 - Bipartite Attention
 - Simplex Attention
 - Duplex Attention
- Testing Results
 - FID Scores
 - Generated Images
- Code Overview
 - Environment Setup
 - Quick-start commands
 - Issues

Input Terms

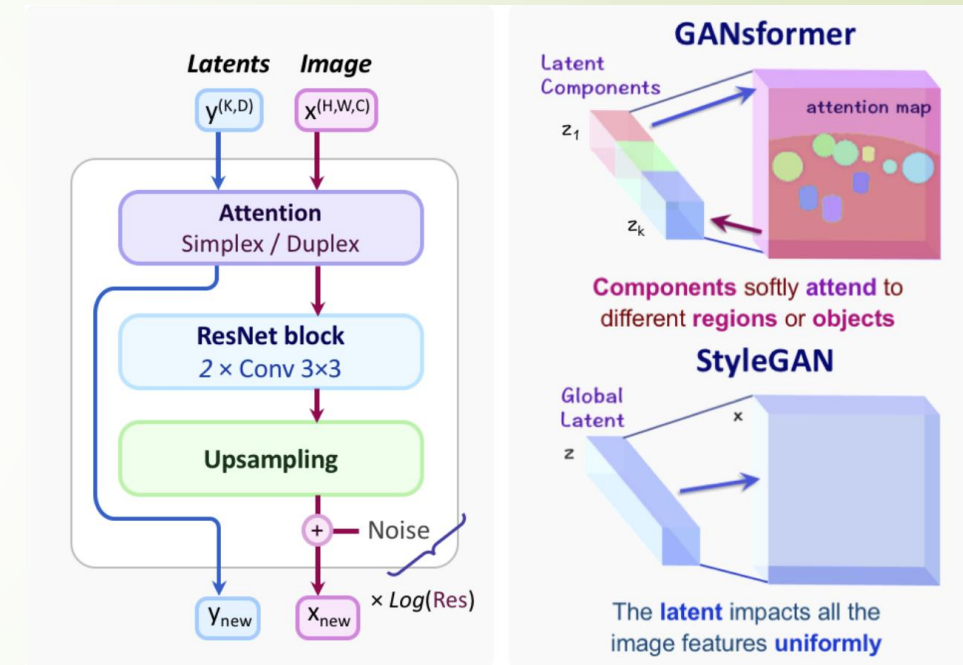
- z : the latent vector
 - Splits into k components $[z_1, \dots, z_k]$
- $y^{m \times d}$: The intermediate latents
 - Akin to output of the mapping network in the StyleGAN architecture
- $x^{n \times d}$: Input Vectors of Dimension d
 - n is *Width \times Height*
 - d is the number of channels



GANsformer Generator Diagram

Input Terms

- z : the latent vector
 - Splits into k components $[z_1, \dots, z_k]$
- $y^{m \times d}$: The intermediate latents
 - Akin to output of the mapping network in the StyleGAN architecture
- $x^{n \times d}$: Input Vectors of Dimension d
 - n is *Width* \times *Height*
 - d is the number of channels
- Notice: inputs in diagram don't match the dimensions of the terms just mentioned
 - Outputs of $q(\cdot)$, $k(\cdot)$, and $v(\cdot)$ transform inputs to their respective dimensions (GANformer only)




GANsformer Generator Diagram

Vanilla Attention

- $Attention(Q, K, V) = softmax\left(\frac{QK^T}{\sqrt{d}}\right) V$
- $a(X) = Attention(q(X), k(X), v(X))$
- $q(\cdot), k(\cdot), v(\cdot)$ map elements to queries, keys, and values
 - Maintain dimensionality
- Poorly scales at higher dimensions
 - $O(n^2 \cdot d)$

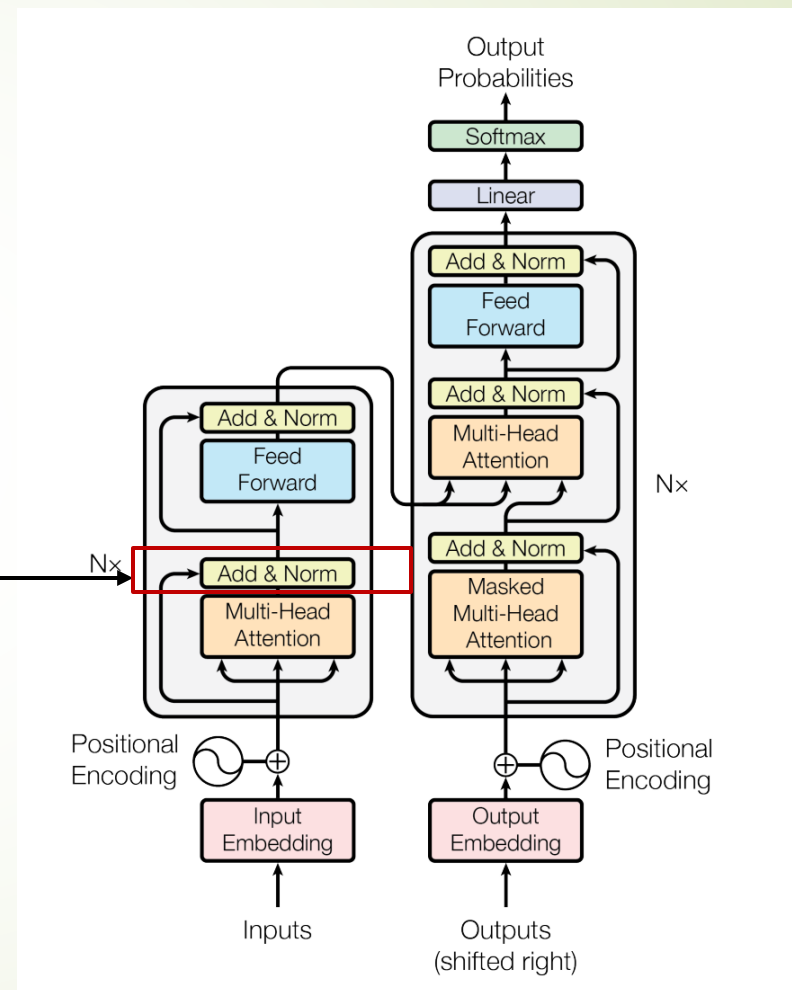
Scaled Dot-Product Attention



Scaled Dot-Product Attention

Vanilla Attention

- $Attention(Q, K, V) = softmax\left(\frac{QK^T}{\sqrt{d}}\right) V$
- $a(X) = Attention(q(X), k(X), v(X))$
- $q(\cdot), k(\cdot), v(\cdot)$ map elements to queries, keys, and values
 - Dimensions are remapped
- Add and LayerNorm right after attention
 - $u^{at}(X) = LayerNorm(X + a(X))$



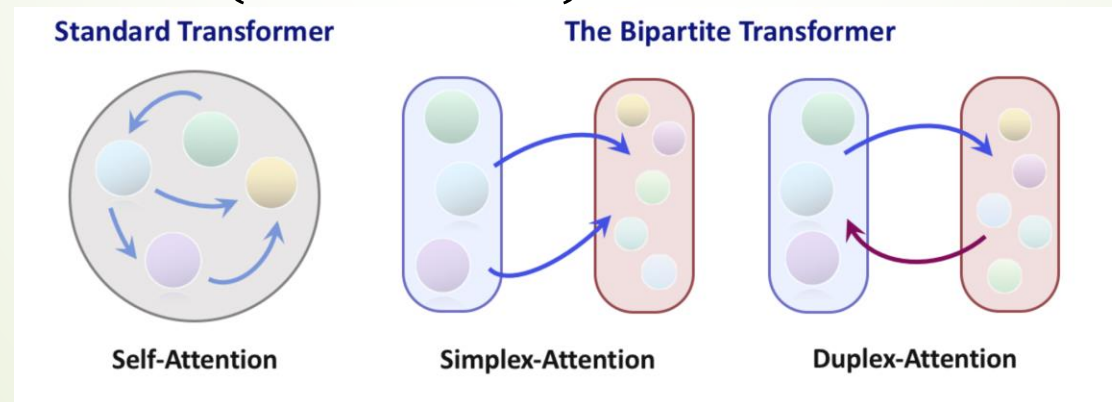
Transformer Architecture

Contents

- Contribution
- Inspirations
- **Attention**
 - Vanilla Attention
 - Bipartite Attention
 - Simplex Attention
 - Duplex Attention
- Testing Results
 - FID Scores
 - Generated Images
- Code Overview
 - Environment Setup
 - Quick-start commands
 - Issues

Bipartite Attention


- $Attention(Q, K, V) = softmax\left(\frac{QK^T}{\sqrt{d}}\right) V$
- $a(X, Y) = Attention(q(X), k(Y), v(Y))$
 - Keys and values are derived from Y (the latents)
- $u^a(X, Y) = LayerNorm(X + a(X, Y))$



Bipartite Attention Efficiency


- Bilinearly Efficient: $O(mn)$
- Recall: $Attention(Q, K, V) = softmax\left(\frac{QK^T}{\sqrt{d}}\right)V$
 - Self attention: $O(n^2 \cdot d)$
 - Matrix multiplications dominate the calculation
- Matrix Multiplication for QK^T
s. t. $Q^{n \times d}, K^{m \times d}: O(ndm) \rightarrow O(mn)$

Standard Transformer




Self-Attention

The Bipartite Transformer



Simplex-Attention



Duplex-Attention

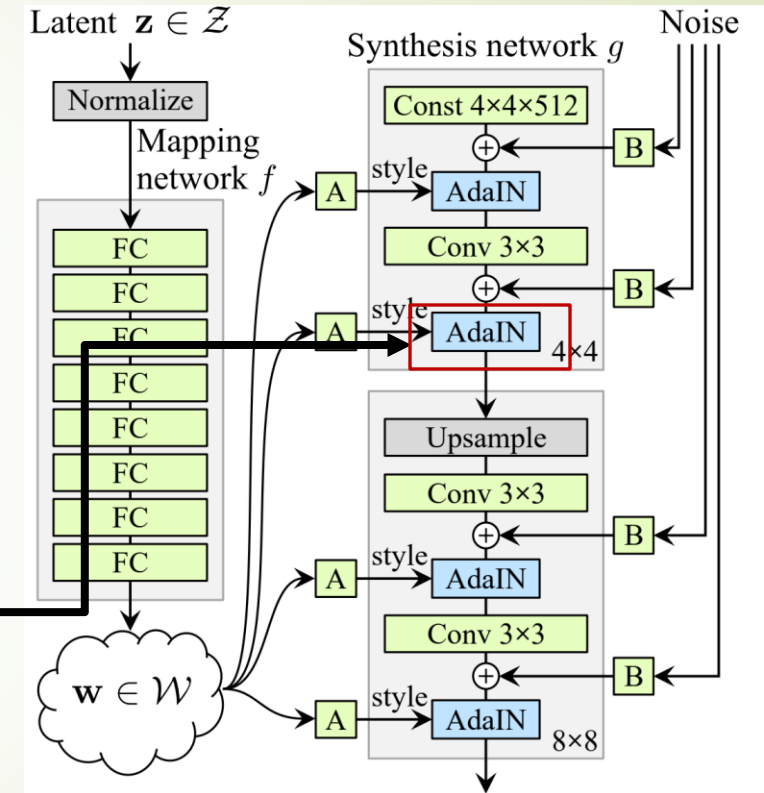
Contents

- Contribution
- Inspirations
- **Attention**
 - Vanilla Attention
 - Bipartite Attention
 - Simplex Attention
 - Duplex Attention
- Testing Results
 - FID Scores
 - Generated Images
- Code Overview
 - Environment Setup
 - Quick-start commands
 - Issues

Simplex Attention

- $Attention(Q, K, V) = softmax\left(\frac{QK^T}{\sqrt{d}}\right) V$
- $a(X, Y) = Attention(q(X), k(Y), v(Y))$
 - Keys and values are derived from Y (the latents)
- Modified AdaIN After
 - $u^s(X, Y) = \gamma(a(X, Y)) \odot \omega(X) + \beta(a(X, Y))$
 - $\omega(X) = \frac{X - \mu(X)}{\sigma(X)}$
 - Operation taken from StyleGAN architecture

Element-wise
multiplication

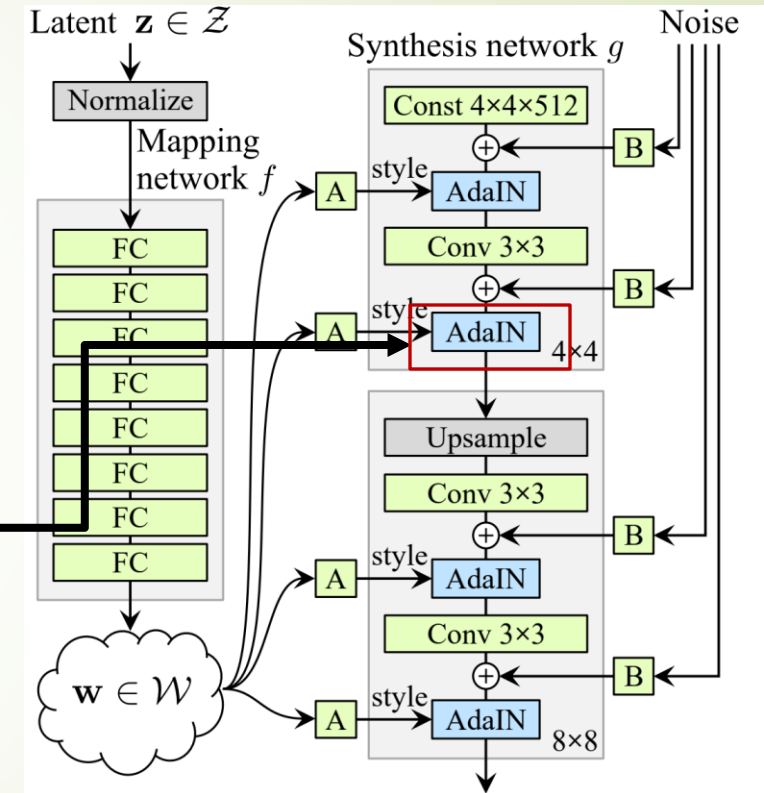


StyleGAN Diagram

Simplex Attention

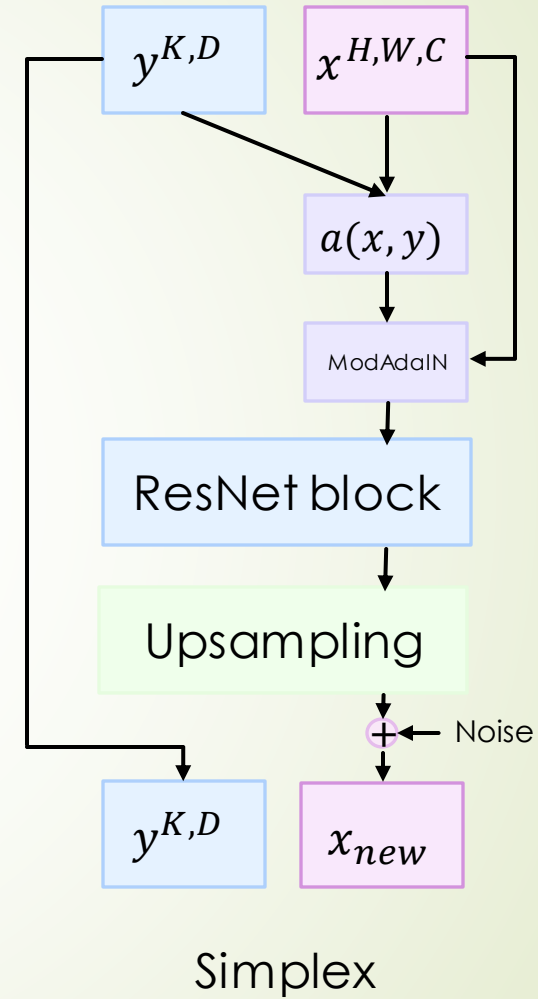

- $Attention(Q, K, V) = softmax\left(\frac{QK^T}{\sqrt{d}}\right) V$
- $a(X, Y) = Attention(q(X), k(Y), v(Y))$
 - Keys and values are derived from Y (the latents)
- Modified AdaIN After
 - $u^s(X, Y) = \gamma(a(X, Y)) \odot \omega(X) + \beta(a(X, Y))$

Element-wise multiplication
 - $\omega(X) = \frac{X - \mu(X)}{\sigma(X)}$
 - Corresponds to StyleGAN architecture
- Latents are global in the generator network



StyleGAN Diagram

Simplex Attention

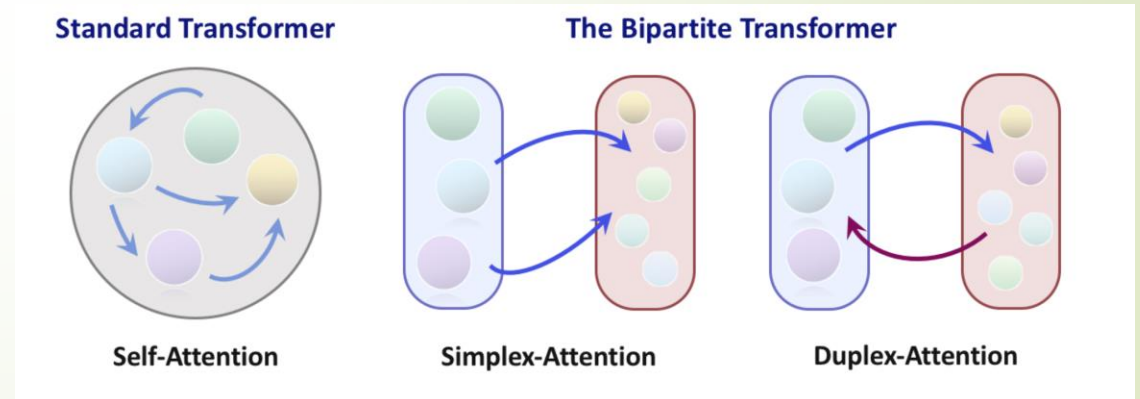


Contents

- Contribution
- Inspirations
- **Attention**
 - Vanilla Attention
 - Bipartite Attention
 - Simplex Attention
 - Duplex Attention
- Testing Results
 - FID Scores
 - Generated Images
- Code Overview
 - Environment Setup
 - Quick-start commands
 - Issues

Duplex Attention


- Recall: $u^a(X, Y) = \text{LayerNorm}(X + a(X, Y))$
- $u^d(X, Y) = \gamma(A(X, K, V)) \odot \omega(X) + \beta(A(X, K, V))$
 - $K = a(Y, X)$
 - Analogous to K-Means
 - Contrast from $u^s(X, Y) = \gamma(a(X, Y)) \odot \omega(X) + \beta(a(X, Y))$



Duplex Attention


- Recall: $u^a(X, Y) = \text{LayerNorm}(X + a(X, Y))$
- $u^d(X, Y) = \gamma(A(X, K, V)) \odot \omega(X) + \beta(A(X, K, V))$
 - $K = a(Y, X)$
 - Analogous to K-Means
 - Contrast from $u^s(X, Y) = \gamma(a(X, Y)) \odot \omega(X) + \beta(a(X, Y))$
- Compute Duplex Attention
 - $Y := u^a(Y, X)$
 - $X := u^d(X, Y)$

Standard Transformer




Self-Attention

The Bipartite Transformer

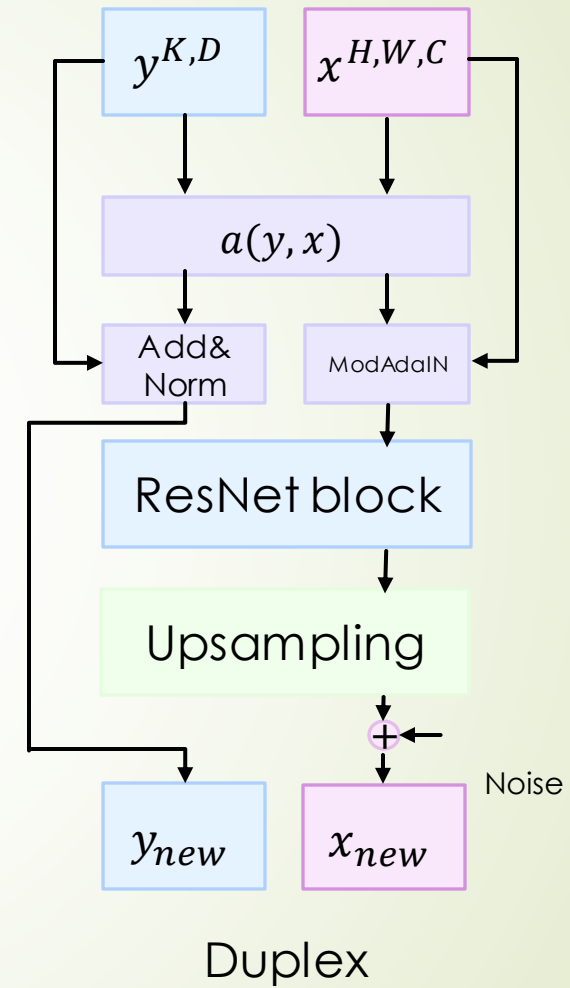
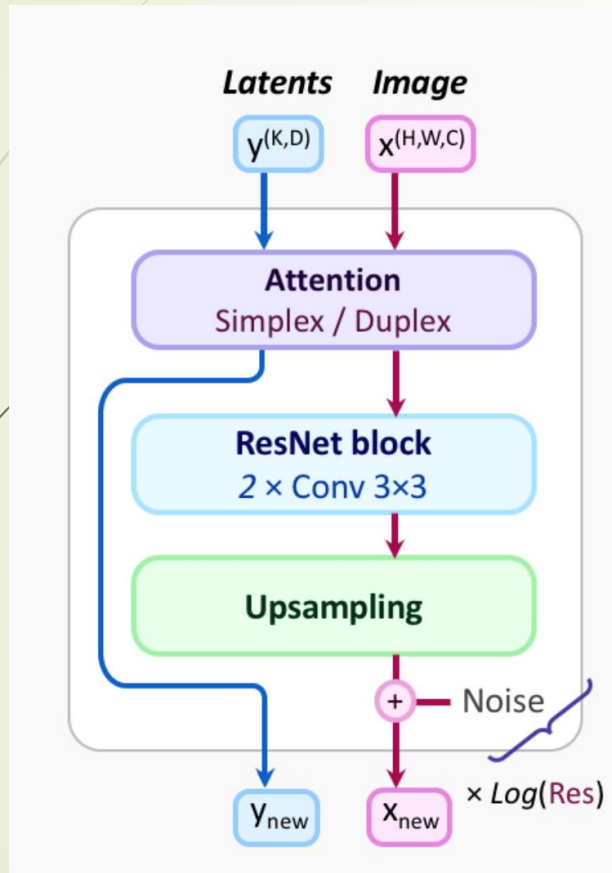


Simplex-Attention

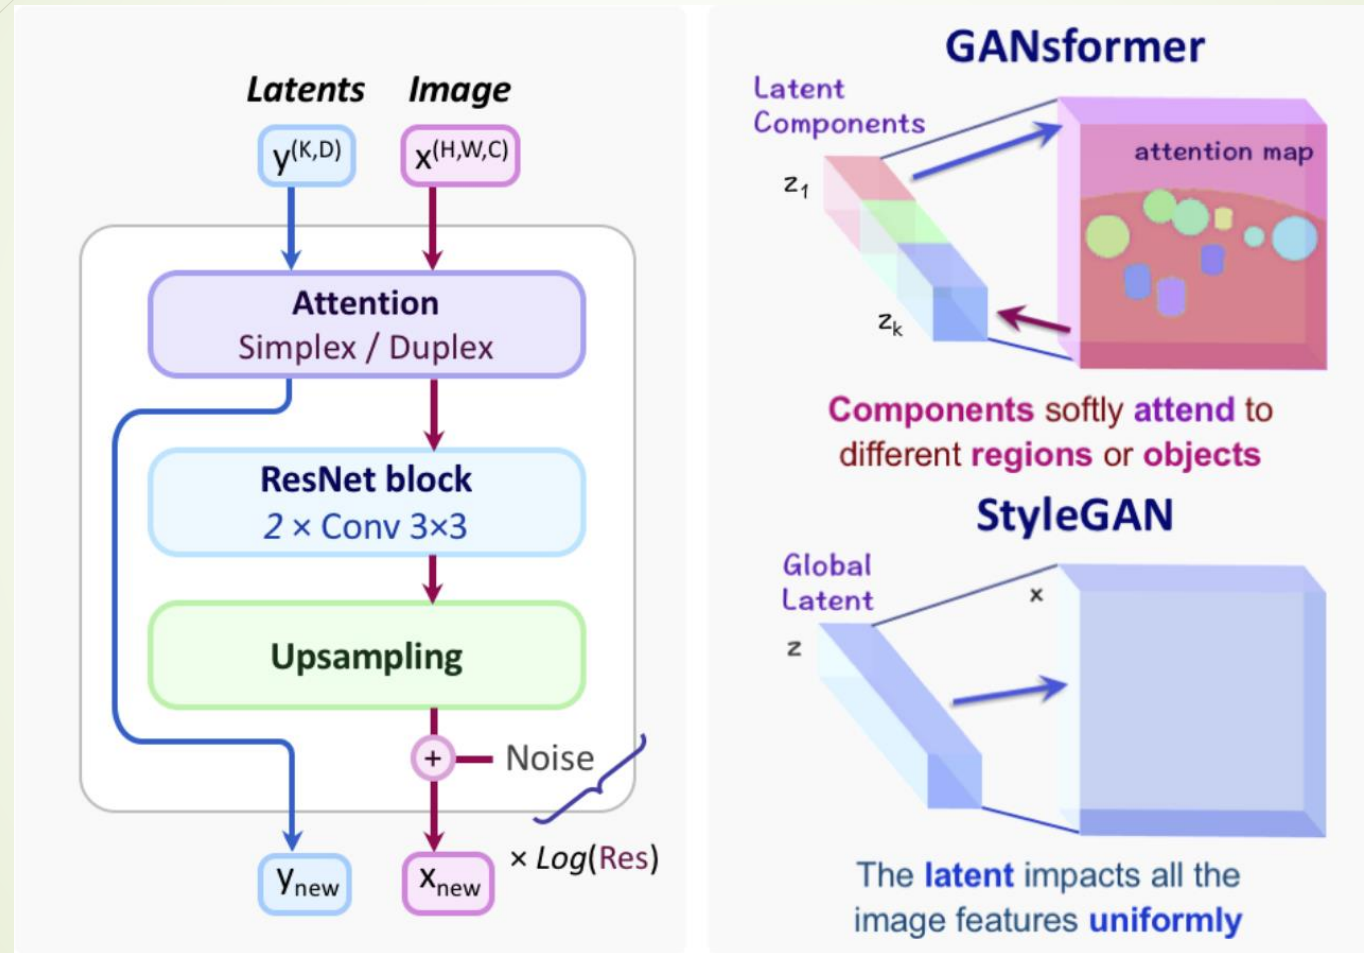


Duplex-Attention

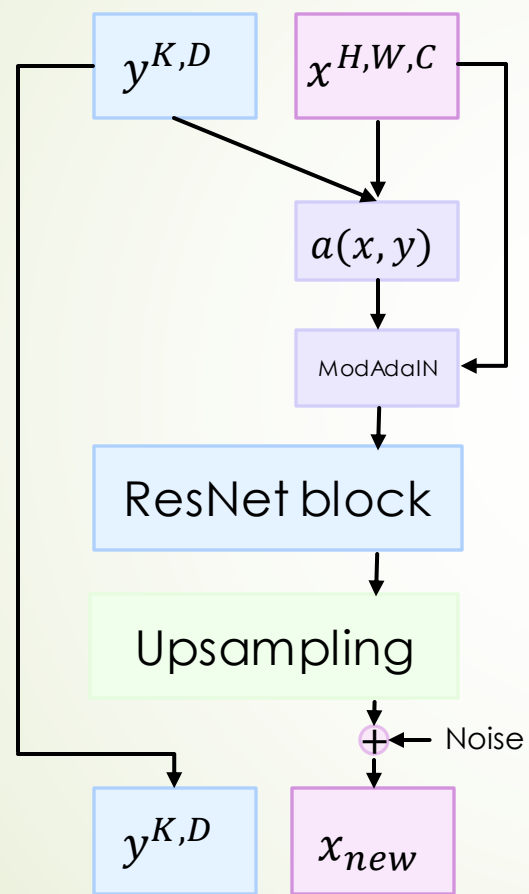
Duplex Attention



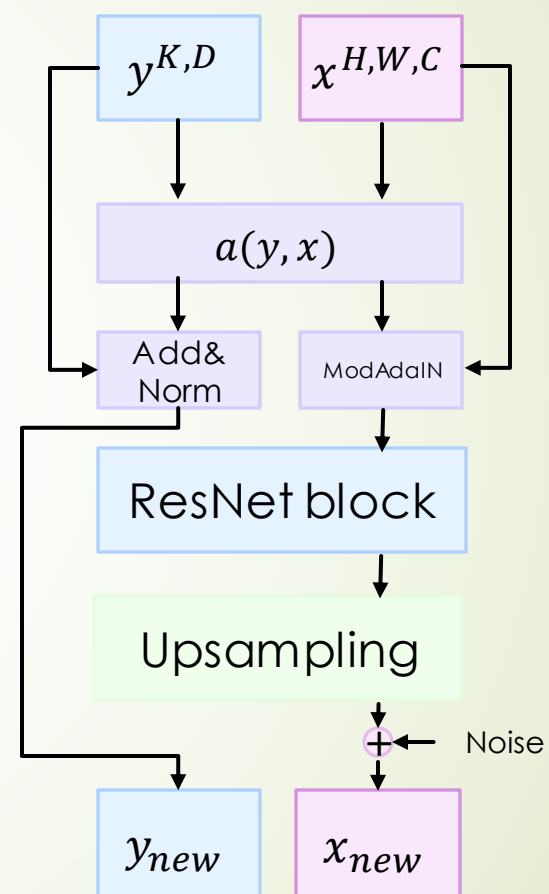
GANsformer Generator Block



Recap




Simplex



Duplex

Generator and Discriminator

- Generator
 - Utilizes attention to allow components to style different regions of the image
- Discriminator
 - Attention applied after every convolution
 - Uses trained embeddings for Y



Contents

- Contribution
- Inspirations
- Attention
 - Vanilla Attention
 - Bipartite Attention
 - Simplex Attention
 - Duplex Attention
- **Testing Results**
 - FID Scores
 - Generated Images
- **Code Overview**
 - Environment Setup
 - Quick-start commands
 - Issues

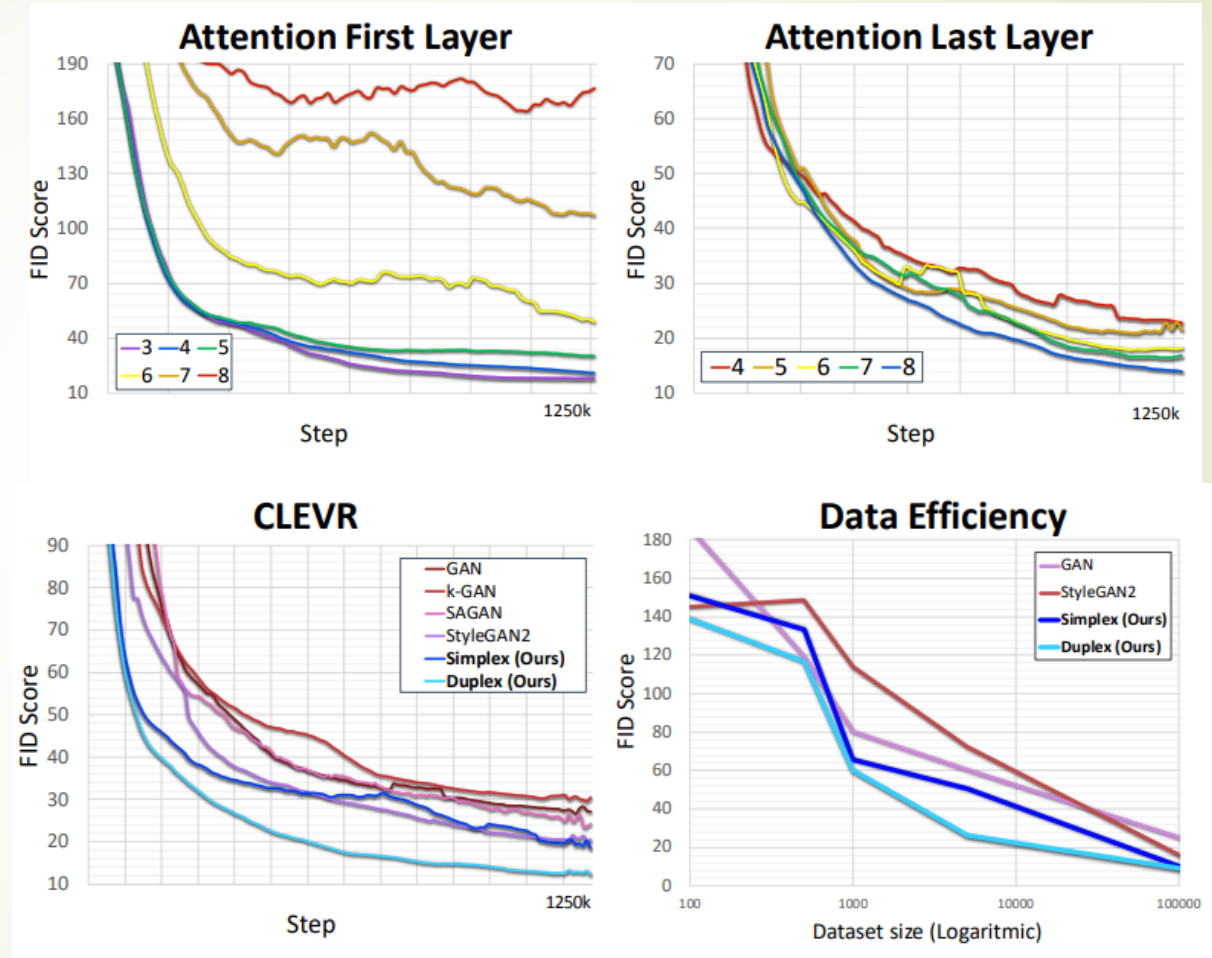
Results

CLEVR					LSUN-Bedroom			
Model	FID ↓	IS ↑	Precision ↑	Recall ↑	FID ↓	IS ↑	Precision ↑	Recall ↑
GAN	25.02	2.17	21.77	16.76	12.16	2.66	52.17	13.63
k-GAN	28.29	2.21	22.93	18.43	69.90	2.41	28.71	3.45
SAGAN	26.04	2.17	30.09	15.16	14.06	2.70	54.82	7.26
StyleGAN2	16.05	2.15	28.41	23.22	11.53	2.79	51.69	19.42
VQGAN	32.60	2.03	46.55	63.33	59.63	1.93	55.24	28.00
GANformer_s	10.26	2.46	38.47	37.76	8.56	2.69	55.52	22.89
GANformer_d	9.17	2.36	47.55	66.63	6.51	2.67	57.41	29.71

FFHQ					Cityscapes			
Model	FID ↓	IS ↑	Precision ↑	Recall ↑	FID ↓	IS ↑	Precision ↑	Recall ↑
GAN	13.18	4.30	67.15	17.64	11.57	1.63	61.09	15.30
k-GAN	61.14	4.00	50.51	0.49	51.08	1.66	18.80	1.73
SAGAN	16.21	4.26	64.84	12.26	12.81	1.68	43.48	7.97
StyleGAN2	9.24	4.33	68.61	25.45	8.35	1.70	59.35	27.82
VQGAN	63.12	2.23	67.01	29.67	173.80	2.82	30.74	43.00
GANformer_s	8.12	4.46	68.94	10.14	14.23	1.67	64.12	2.03
GANformer_d	7.42	4.41	68.77	5.76	5.76	1.69	48.06	33.65

Results

- Achieves better FID scores in fewer training steps than previous models




Contents

- Contribution
- Inspirations
- Attention
 - Vanilla Attention
 - Bipartite Attention
 - Simplex Attention
 - Duplex Attention
- **Testing Results**
 - FID Scores
 - Generated Images
- **Code Overview**
 - Environment Setup
 - Quick-start commands
 - Issues

Our Evaluations

▀ ffhq-snapshot FID: 6.1067




▀ cityscapes-snapshot FID: 7.8310



Baseline Model for Cityscapes-2048







FFHQ-1024 Baseline: Good



FFHQ-1024 Baseline: Bad



Attention Maps

- Each color corresponds to a different component of the k latents that influence that region



Contents

- Contribution
- Inspirations
- Attention
 - Vanilla Attention
 - Bipartite Attention
 - Simplex Attention
 - Duplex Attention
- Testing Results
 - FID Scores
 - Generated Images
- **Code Overview**
 - Environment Setup
 - Quick-start commands
 - Issues

Running the Code On Windows Environment Setup

- Clone from <https://github.com/dorarad/gansformer>
- Create a Python 3.7 virtual environment and install packages from the requirements.txt file
- Install Visual Studio 2017 and the MSVC compiler
 - Edit ./dnnlib/tflib/custom_ops.py and add your compiler path to the **compiler_bindir_search_path** variable

```
22 compiler_bindir_search_path = [  
23     "C:/Program Files (x86)/Microsoft Visual Studio/2017/Community/VC/Tools/MSVC/14.14.26428/bin/Hostx64/x64",  
24     #"C:/Program Files (x86)/Microsoft Visual Studio/2019/Community/VC/Tools/MSVC/14.23.28105/bin/Hostx64/x64",  
25     "C:/Program Files (x86)/Microsoft Visual Studio/2017/Community/VC/Tools/MSVC/14.16.27023/bin/Hostx64/x64",  
26     "C:/Program Files (x86)/Microsoft Visual Studio 14.0/vc/bin",  
27 ]
```


Running the Code On Windows Environment Setup

- Clone from <https://github.com/dorarad/gansformer>
- Create a Python 3.7 virtual environment and install packages from the requirements.txt file
- Install Visual Studio 2017 and the MSVC compiler
 - Edit ./dnnlib/tflib/custom_ops.py and add your compiler path to the **compiler_bindir_search_path** variable
- Download and install CUDA 10.0
- Download cuDNN v7.6.5 (November 5th, 2019), for CUDA 10.0
 - Follow the instructions in section 3.3 in the following link
 - <https://docs.nvidia.com/deeplearning/cudnn/install-guide/index.html#installwindows>

Contents

- Contribution
- Inspirations
- Attention
 - Vanilla Attention
 - Bipartite Attention
 - Simplex Attention
 - Duplex Attention
- Testing Results
 - FID Scores
 - Generated Images
- **Code Overview**
 - Environment Setup
 - Quick-start commands
 - Issues

Running the Code On Windows

- `python generate.py --gpus 0 --model gdrive:<dataset>-snapshot.pkl --output-dir images --images-num 32`
 - Replace “<dataset>” with different supported dataset name
- `python run_network.py --train --gpus 0 --ganformer-default --expname <dataset>-pretrained --dataset <dataset> --pretrained-pkl gdrive:<dataset>-snapshot.pkl`
 - Trains a pre-existing network
- `Python run_network.py --help`
 - Help menu for more options

Contents

- Contribution
- Inspirations
- Attention
 - Vanilla Attention
 - Bipartite Attention
 - Simplex Attention
 - Duplex Attention
- Testing Results
 - FID Scores
 - Generated Images
- **Code Overview**
 - Environment Setup
 - Quick-start commands
 - Issues

Issues We Ran Into

- Could not train a model using duplex attention
 - Enabling causes errors
 - Could be an issue with the dependencies
- Often ran out of GPU memory
 - Code designed with 12GB GPUs in mind
 - Led to crashing
 - Couldn't generate attention maps

```
Produce visualizations...
0%|
2022-03-27 11:36:14.279813: E tensorflow/stream_executor/cuda/cuda_driver.cc:828] failed to allocate 4.00G (4291821568 bytes) from device: CUDA_ERROR_OUT_OF_MEMORY: out of memory
2022-03-27 11:36:14.455244: E tensorflow/stream_executor/cuda/cuda_driver.cc:828] failed to allocate 3.60G (3862639360 bytes) from device: CUDA_ERROR_OUT_OF_MEMORY: out of memory
2022-03-27 11:36:14.623482: E tensorflow/stream_executor/cuda/cuda_driver.cc:828] failed to allocate 3.24G (3476375296 bytes) from device: CUDA_ERROR_OUT_OF_MEMORY: out of memory
2022-03-27 11:36:14.788041: E tensorflow/stream_executor/cuda/cuda_driver.cc:828] failed to allocate 2.91G (3128737792 bytes) from device: CUDA_ERROR_OUT_OF_MEMORY: out of memory
100%|#####| 1/1 [00:15<00:00, 15.40s/it]
```


Questions?